

LLM Benchmark Report

2026-04-02 00:24

Configuration

Parameter	Value
Model	sagemaker/jumpstart-dft-hf-reasoning-qwen3-4b-20260323-223535
API	litellm
Tokenizer	Qwen/Qwen3-4B
Modality	text
Mean Input Tokens	1024
Mean Output Tokens	256
Requests / Step	32
Warmup Requests	2
Concurrency Steps	1, 2, 4, 8, 16


Highlights

Peak Throughput	227.1 tok/s @ concurrency 8
Total Errors	0


Concurrency Sweep Summary

Concurrency	Throughput (tok/s)	Req/min	E2E p50 (s)	E2E p99 (s)	TTFT p50 (s)	ITL p50 (s)	Errors
1	70.0	16.5	3.495	4.071	0.1504	0.013703	0
2	120.8	28.5	3.692	5.168	0.1211	0.014462	0
4	201.1	47.5	3.799	4.382	0.1214	0.014884	0
8	227.1	53.7	4.083	10.820	0.1397	0.015890	0
16	206.5	48.9	4.721	7.139	2.3932	0.017971	0

Throughput & Request Rate vs Concurrency



Latency vs Concurrency



Concurrency: 1

Throughput: 70.0 tok/s | Completed: 32 | Requests/min: 16.5 | Errors: 0

Metric	p25	p50	p75	p90	p95	p99	mean	min	max	stddev
ITL (s)	0.0137	0.0137	0.0137	0.0138	0.0143	0.0161	0.0138	0.0136	0.0168	0.000571
TTFT (s)	0.1475	0.1504	0.1567	0.1833	0.3097	0.7619	0.1854	0.1214	0.9037	0.1417
E2E Latency (s)	3.41	3.49	3.61	3.68	3.75	4.07	3.52	3.23	4.19	0.1761
Throughput (tok/s)	72.57	72.79	73.06	73.23	73.26	73.42	72.22	59.49	73.49	2.53
Input Tokens	1018.0	1025.0	1028.2	1038.2	1045.5	1054.5	1024.3	999	1056	11.54
Output Tokens	248.0	253.5	258.8	264.9	266.4	267.0	253.8	236	267	8.04

ITL (s)

TTFT (s)

E2E Latency (s)

Throughput (tok/s)

Input Tokens

Output Tokens

— Inter-Token Latency

— Time to First Token

— End-to-End Latency

— Per-Request Output Throughput

— Input Token Count

— Output Token Count

Concurrency: 2

Throughput: 120.8 tok/s | Completed: 32 | Requests/min: 28.5 | Errors: 0

Metric	p25	p50	p75	p90	p95	p99	mean	min	max	stddev
ITL (s)	0.0144	0.0145	0.0145	0.0146	0.0171	0.0204	0.0148	0.0144	0.0205	0.001437
TTFT (s)	0.1109	0.1211	0.1251	0.1448	0.2902	1.28	0.1784	0.1044	1.64	0.2740
E2E Latency (s)	3.60	3.69	3.80	3.85	4.42	5.17	3.77	3.43	5.20	0.3782
Throughput (tok/s)	68.88	68.99	69.12	69.20	69.30	69.34	67.74	48.71	69.35	4.88
Input Tokens	1018.0	1025.0	1028.2	1038.2	1045.5	1054.5	1024.3	999	1056	11.54
Output Tokens	248.0	253.5	258.8	264.9	266.4	267.0	253.9	236	267	8.04

- ITL (s)

TTFT (s)

E2E Latency (s)

Throughput (tok/s)

Input Tokens

Output Tokens
- Inter-Token Latency

— Time to First Token

— End-to-End Latency

— Per-Request Output Throughput

— Input Token Count

— Output Token Count

Concurrency: 4

Throughput: 201.1 tok/s | Completed: 32 | Requests/min: 47.5 | Errors: 0

Metric	p25	p50	p75	p90	p95	p99	mean	min	max	stddev
ITL (s)	0.0149	0.0149	0.0149	0.0163	0.0165	0.0166	0.0151	0.0148	0.0166	0.000549
TTFT (s)	0.1141	0.1214	0.1297	0.4783	0.5678	0.5946	0.1753	0.1035	0.5974	0.1497
E2E Latency (s)	3.71	3.80	3.93	4.06	4.27	4.38	3.84	3.53	4.42	0.2032
Throughput (tok/s)	66.94	67.03	67.11	67.18	67.24	67.29	66.22	60.13	67.30	2.22
Input Tokens	1018.0	1025.0	1028.2	1038.2	1045.5	1054.5	1024.3	999	1056	11.54
Output Tokens	248.0	254.0	258.8	264.9	266.4	267.0	253.9	236	267	8.11

- ITL (s)

TTFT (s)

E2E Latency (s)

Throughput (tok/s)

Input Tokens

Output Tokens
- Inter-Token Latency

— Time to First Token

— End-to-End Latency

— Per-Request Output Throughput

— Input Token Count

— Output Token Count

Concurrency: 8

Throughput: 227.1 tok/s | Completed: 32 | Requests/min: 53.7 | Errors: 0

Metric	p25	p50	p75	p90	p95	p99	mean	min	max	stddev
ITL (s)	0.0158	0.0159	0.0163	0.0180	0.0186	0.0419	0.0174	0.0157	0.0523	0.006436
TTFT (s)	0.1243	0.1397	0.6903	1.05	1.39	7.36	0.6881	0.1069	9.86	1.72
E2E Latency (s)	4.00	4.08	4.29	4.50	4.72	10.82	4.43	3.72	13.52	1.68
Throughput (tok/s)	61.11	62.78	63.20	63.57	63.61	63.67	60.03	19.08	63.69	8.12
Input Tokens	1018.8	1025.0	1027.2	1030.9	1044.4	1054.5	1024.2	999	1056	10.90
Output Tokens	248.0	253.5	258.0	264.0	266.4	267.0	253.5	236	267	7.98

- ITL (s)

TTFT (s)

E2E Latency (s)

Throughput (tok/s)

Input Tokens

Output Tokens
- Inter-Token Latency

— Time to First Token

— End-to-End Latency

— Per-Request Output Throughput

— Input Token Count

— Output Token Count

Concurrency: 16

Throughput: 206.5 tok/s | Completed: 32 | Requests/min: 48.9 | Errors: 0

Metric	p25	p50	p75	p90	p95	p99	mean	min	max	stddev
ITL (s)	0.0175	0.0180	0.0216	0.0233	0.0244	0.0285	0.0199	0.0173	0.0302	0.003054
TTFT (s)	1.57	2.39	3.14	4.09	4.50	4.93	2.50	1.04	5.03	1.10
E2E Latency (s)	4.50	4.72	5.42	6.01	6.17	7.14	5.06	4.24	7.52	0.7434
Throughput (tok/s)	46.41	55.41	56.77	57.29	57.34	57.47	50.98	33.00	57.52	6.79
Input Tokens	1018.8	1025.0	1027.2	1030.9	1045.5	1054.5	1024.1	999	1056	11.20
Output Tokens	248.0	252.5	258.8	265.0	267.0	267.0	253.4	236	267	8.73

ITL (s)

TTFT (s)

E2E Latency (s)

Throughput (tok/s)

Input Tokens

Output Tokens

— Inter-Token Latency

— Time to First Token

— End-to-End Latency

— Per-Request Output Throughput

— Input Token Count

— Output Token Count